

Debugging

Joseph Hallett

November 4, 2024



Whats all this about?

Writing programs is hard

- ▶ We should have *strategies* and *tools* for when things go wrong

Lets point you towards some!

An example program

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    char message[128];
    size_t message_len = 256;
    char timestamp[128];
    time_t t;
    struct tm *tmp;
    FILE *file = fopen(argv[1], "a+");

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}
```

Lets compile!

```
make journal
```

```
cc  journal.c -o journal
```

And when we run...

```
./journal <<<"Hello World!"
```

```
Segmentation fault (core dumped)
```

Segfaults are annoying

Programs randomly crashing is extremely irritating

- ▶ ...but a fact of life
- ▶ ... *we are* programmers
- ▶ ... it is *our* job to fix them

Lets discuss some tools for spotting errors

AddressSanitizer Adds extra debugging checks

GDB the GNU Debugger

Strace Systemcall tracer

Ltrace Library tracer

Valgrind Memory error detector

AddressSanitizer

Tool from the LLVM project (clang) to give more information about crashes **at runtime**.
Other sanitizers exist:

UndefinedBehaviorSanitizer spots bad compiler practice

ThreadSanitizer spots bad multithreading practice

LeakSanitizer spots bad memory management

To use, compile with clang and add the `-fsanitize=address` flag

```
clang journal.c -fsanitize=address -o journal
```

Downsides

- ▶ Your program will use more memory and be slower
- ▶ Your program may be more easy to hack ;-)

Lets go!

Looks like its a NULL pointer dereference?

```
$ ./journal <<<"hello"
AddressSanitizer:DEADLYSIGNAL
=====
==302530==ERROR: AddressSanitizer: SEGV on unknown address 0x0000000000c0 (pc
==302530==The signal is caused by a READ memory access.
==302530==Hint: address points to the zero page.
#0 0x7f99bd58e470 (/lib64/libc.so.6+0x66470) (BuildId: 37e4ac6a7fb96950b)
#1 0x453325 (/home/jh18636/journal+0x453325) (BuildId: becf911c680ef0557)
#2 0x5115f0 (/home/jh18636/journal+0x5115f0) (BuildId: becf911c680ef0557)
#3 0x7f99bd5627e4 (/lib64/libc.so.6+0x3a7e4) (BuildId: 37e4ac6a7fb96950b)
#4 0x42975d (/home/jh18636/journal+0x42975d) (BuildId: becf911c680ef0557)

AddressSanitizer can not provide additional info.
SUMMARY: AddressSanitizer: SEGV (/lib64/libc.so.6+0x66470) (BuildId: 37e4ac6a
==302530==ABORTING
```


GDB: The GNU Debugger

Horribly powerful

- ▶ Can step through a program at an assembly level
- ▶ Can watch registers and the stack change line by line
- ▶ Can program scripts and to react when certain things happen
- ▶ Can debug systems remotely
- ▶ Ported to every sort of computer you could wish for

Horribly unintuitive

- ▶ It has a built in GUI but you'll wish it didn't (layout regs)
- ▶ Lots of unintuitive single letter commands
- ▶ Cryptic output
- ▶ But its the standard debugger!
 - ▶ Others (and GUIs) exist, but still worth having a basic knowledge of the CLI tool

Okay, lets try and debug

- ▶ run runs your program (with arguments if passed)
- ▶ bt gives you a backtrace by reading out the stack

```
$ gdb ./journal
Reading symbols from ./journal...
(No debugging symbols found in ./journal)
(gdb) run <<<"hello"
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal <<<"hello"
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib64/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
__vfprintf_internal (s=0x0, format=0x402026 "%s: %s\n", ap=ap@entry=0x7fffffffde50, mode_flags=mode_
722    ORIENT;
(gdb) bt
#0 __vfprintf_internal (s=0x0, format=0x402026 "%s: %s\n",
    ap=ap@entry=0x7fffffffde50, mode_flags=mode_flags@entry=0)
    at vfprintf-internal.c:722
#1 0x00007ffff7e2360a in __fprintf (stream=<optimized out>,
    format=<optimized out>) at fprintf.c:32
#2 0x000000000040125f in main ()
```

Lets make it a *little* easier

- ▶ -g adds debugging informations.
- ▶ -Og optimizes for debuggability

```
$ cc -Og -g journal.c -o journal
$ gdb ./journal
(gdb) run <<<"hello"
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal <<<"hello"
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib64/libthread_db.so.1".

Program received signal SIGSEGV, Segmentation fault.
__memcpy_avx_unaligned_erms () at ../sysdeps/x86_64/multiarch/memmove-vec-unaligned-erms.S:333
Downloading 0.01 MB source file /usr/src/debug/glibc-2.36.9000-19.fc38.x86_64/string/./sysdeps/x86_64/333
333      movl    %ecx, -4(%rdi, %rdx)
(gdb) bt
#0  __memcpy_avx_unaligned_erms ()
    at ../sysdeps/x86_64/multiarch/memmove-vec-unaligned-erms.S:333
#1  0x00007ffff7e496ac in __GI__getdelim (
    lineptr=lineptr@entry=0x7ffffffffffdf0, n=n@entry=0x7ffffffffffdf8,
    delimiter=delimiter@entry=10, fp=0x7ffff7fa5aa0 <_IO_2_1_stdin_>)
    at iogetdelim.c:111
#2  0x00007ffff7e237d1 in __getline (lineptr=lineptr@entry=0x7ffffffffffdf0,
    n=n@entry=0x7ffffffffffdf8, stream=<optimized out>) at getline.c:28
#3  0x00000000004011d6 in main (argc=<optimized out>, argv=<optimized out>)
    at journal.c:14
```

Looks like it all went wrong on line 14 of journal.c...

Breakpoints let us stop when we reach a line of the program or an address.

- ▶ Create them with `b address`
- ▶ Delete them with `d`
- ▶ Run on to the next one with `c`

```
(gdb) b journal.c:14
Breakpoint 2 at 0x4011ba: file journal.c, line 14.
(gdb) run <<<"hello"
The program being debugged has been started already.
Start it from the beginning? (y or n) y
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal <<<"hello"
[{}Thread debugging using libthread\textsubscript{db} enabled]
Using host libthread\textsubscript{db} library "/lib64/libthread\textsubscript{db.so.1}".

Breakpoint 2, main (argc=<optimized out>, argv=<optimized out>) at journal.c:14
14  getline(&message, &message\textsubscript{len}, stdin);
(gdb) inspect message
\${3} = "@$\backslashbackslash$000$\backslashbackslash$000$\backslashbackslash$000$\backslashbackslash$000$\backslashbackslash$000$\backslashbackslash$000"
(gdb) inspect message\textsubscript{len}
\${4} = 256
(gdb) d
Delete all breakpoints? (y or n) y
(gdb)
```

If in doubt... read the manual

In man 3 getline:

*getline() reads an entire line from stream, storing the address of the buffer containing the text into *lineptr. The buffer is null-terminated and includes the newline character, if one was found.*

*If *lineptr is set to NULL before the call, then getline() will allocate a buffer for storing the line. This buffer should be freed by the user program even if getline() failed.*

*Alternatively, before calling getline(), *lineptr can contain a pointer to a malloc(3)-allocated buffer *n bytes in size. If the buffer is not large enough to hold the line, getline() resizes it with realloc(3), updating *lineptr and *n as necessary.*

Well we're passing a statically allocated buffer... lets fix that.

A new example program

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
    size_t message_len;
    char timestamp[128];
    time_t t;
    struct tm *tmp;
    FILE *file = fopen(argv[1], "a+");

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}
```

```
cc -g -Og journal2.c -o journal2
```

And now when we run...

```
$ ./journal2 <<<"hello"
Segmentation fault (core dumped)

$ gdb ./journal2
(gdb) run <<<"hello"
Starting program: /home/joseph/Repos/Talks/COMS10012-Software-Tools/Debugging/journal2 <<<"hello"

Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7e2de82 in __vfprintf_internal () from /lib64/libc.so.6
Missing separate debuginfos, use: dnf debuginfo-install glibc-2.36.9000-19.fc38.x86_64
(gdb) bt
#0 0x00007ffff7e2de82 in __vfprintf_internal () from /lib64/libc.so.6
#1 0x00007ffff7e2360a in fprintf () from /lib64/libc.so.6
#2 0x0000000000401225 in main (argc=<optimized out>, argv=<optimized out>) at journal2.c:20
(gdb)
```

...well, we got further...

We could continue with gdb

GDB is an extremely powerful debugging tool

- ▶ Its also *really* hard to use
- ▶ See *Computer Systems B* next year, or *Systems and Software Security* at Masters level
- ▶ If you're on a Mac or BSD box check out `lldb`
- ▶ Or for a proper tutorial the documentation it refers you to *every time you open it*.

It is *well worth your time to learn...*

- ▶ But *this course* is about *Software Tools* and I want to show you *more* of them

`<<input` runs your program with input

`b` set breakpoints

`c` continue after hitting a breakpoint

`bt` get a backtrace

`info` get information about registers or variables or anything else

`x` examine a variable/pointer

`disas` see the assembly code you're running

`help` get help!

Strace

The strace tool lets you trace what systemcalls a program uses

- ▶ On OpenBSD see ktrace and kdump
- ▶ On MacOS/FreeBSD see dtruss and dtrace

Lets run it!

```
make journal2
strace ./journal2 <<<'Hello' 2>&1
```

```
execve("./journal2", [".journal2"], 0x7fffe3c9beb0 /* 23 vars */) = 0
brk(NULL) = 0x56769b94b000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
fstat(3, {st_mode=S_IFREG|0644, st_size=311295, ...}) = 0
mmap(NULL, 311295, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7c742720d000
close(3) = 0
openat(AT_FDCWD, "/usr/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0\220^\2\0\0\0\0\0"..., 832) = 832
pread64(3, "\6\0\0\0\4\0\0\0@\0\0\0\0\0\0@\0\0\0\0\0\0@\0\0\0\0\0\0"..., 784, 64) = 784
fstat(3, {st_mode=S_IFREG|0755, st_size=1948952, ...}) = 0
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7c742720b000
pread64(3, "\6\0\0\0\4\0\0\0@\0\0\0\0\0\0@\0\0\0\0\0\0@\0\0\0\0\0\0"..., 784, 64) = 784
mmap(NULL, 1973104, PROT_READ, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7c7427029000
mmap(0x7c742704d000, 1421312, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x24000) = 0x7c742704d000
mmap(0x7c74271a8000, 348160, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x17f000) = 0x7c74271a8000
mmap(0x7c74271fd000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1d3000) = 0x7c74271fd000
mmap(0x7c7427203000, 31600, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x7c7427203000
close(3) = 0
mmap(NULL, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7c7427026000
arch_prctl(ARCH_SET_FS, 0x7c7427026740) = 0
set_tid_address(0x7c7427026a10) = 32033
set_robust_list(0x7c7427026a20, 24) = 0
rseq(0x7c7427027060, 0x20, 0, 0x53053053) = 0
mprotect(0x7c74271fd000, 16384, PROT_READ) = 0
```

Too much output!

strace lets you use regexp to filter what syscalls you look at

- ▶ ...or you could just use grep...

```
make journal2
strace -e '/open.*' ./journal2 <<<hello 2>&1
```

```
cc  journal2.c -o journal2
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, NULL, O_RDWR|O_CREAT|O_APPEND, 0666) = -1 EFAULT (Bad address)
openat(AT_FDCWD, "/etc/localtime", O_RDONLY|O_CLOEXEC) = 3
--- SIGSEGV {si_signo=SIGSEGV, si_code=SEGV_MAPERR, si_addr=0xc0} ---
+++ killed by SIGSEGV (core dumped) +++
```

Or if you prefer OpenBSD

```
ktrace ./journal2 <<<hello  
kdump | grep -A 2 -B 2 open
```

```
--  
63085 journal2 CALL close(3)  
63085 journal2 RET close 0  
63085 journal2 CALL open(0x34d8c02f033,0x10000<O_RDONLY|O_CLOEXEC>)  
63085 journal2 NAMI "/usr/lib/libc.so.100.3"  
63085 journal2 RET open 3  
63085 journal2 CALL fstat(3,0x70bf220f2088)  
63085 journal2 STRU struct stat { dev=1077, ino=3784328, mode=-r--r--r-- , nlink=1, uid=0<"root">,  
--  
63085 journal2 CALL mprotect(0x34debfa5000,0x1000,0x1<PROT_READ>)  
63085 journal2 RET mprotect 0  
63085 journal2 CALL open(0,0x20a<O_RDWR|O_APPEND|O_CREAT>,0666<S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_I  
63085 journal2 RET open -1 errno 14 Bad address  
63085 journal2 CALL kbind(0x70bf220f2268,24,0x5e8161d98e625637)  
63085 journal2 RET kbind 0  
--
```

Oh yeah... we forgot an arg

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
    size_t message_len;
    char timestamp[128];
    time_t t;
    struct tm *tmp;
    FILE *file = fopen(argv[1], "a+");

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}
```

Lets fix that...

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
    size_t message_len;
    char timestamp[128];
    time_t t;
    struct tm *tmp;

    if (argc < 2) { printf("Usage: %s path/to/log\n", argv[0]); exit(1); };
    FILE *file = fopen(argv[1], "a+");

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}
```

And if you cant spot the difference

```
diff -u journal{2,3}.c
```

```
--- journal2.c 2024-02-07 11:14:29.060025998 +0000
+++ journal3.c 2024-02-07 12:16:09.220079001 +0000
@@ -8,6 +8,8 @@
     char timestamp[128];
     time_t t;
     struct tm *tmp;
+
+ if (argc < 2) { printf("Usage: %s path/to/log\n", argv[0]); exit(1); };
     FILE *file = fopen(argv[1], "a+");

     printf("Type your log: ");
```

Now when we run!

```
$ ./journal3 documents/log.txt <<<hello  
Segmentation fault (core dumped)
```

Lets try `ltrace` this time (no equivalent on other platforms)...

- It traces *library* calls

ltrace and a bit more strace

```
make journal3
ltrace ./journal3 documents/log.txt <<<hello 2>&1
```

```
fopen("documents/log.txt", "a+")      = 0
printf("Type your log: ")             = 15
getline(0x7fffd196b0018, 0x7fffd196b0020, 0x76a22f8538e0, 0x7fffd196b0020) = 6
time(0)                               = 1707308599
localtime(0x7fffd196b0028)             = 0x76a22f85a320
strftime("20", 256, "%C", 0x76a22f85a320) = 2
fprintf(0, "%s: %s\n", "20", "hello\n" <no return ...>
--- SIGSEGV (Segmentation fault) ---
+++ killed by SIGSEGV +++
```

```
strace -e openat ./journal3 documents/log.txt <<<hello 2>&1
```

```
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "/usr/lib/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
openat(AT_FDCWD, "documents/log.txt", O_RDWR|O_CREAT|O_APPEND, 0666) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/localtime", O_RDONLY|O_CLOEXEC) = 3
--- SIGSEGV {si_signo=SIGSEGV, si_code=SEGV_MAPERR, si_addr=0xc0} ---
+++ killed by SIGSEGV (core dumped) ++
```

Lets fix that...

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <errno.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
    size_t message_len;
    char timestamp[128];
    time_t t;
    struct tm *tmp;

    if (argc < 2) { printf("Usage: %s path/to/log\n", argv[0]); exit(1); };
    FILE *file = fopen(argv[1], "a+");
    if (file == NULL) {
        perror("Failed to open log");
        exit(2);
    }

    printf("Type your log: ");
    getline(&message, &message_len, stdin);

    t = time(NULL);
    tmp = localtime(&t);
    strftime(timestamp, 256, "%C", tmp);

    fprintf(file, "%s: %s\n", timestamp, message);
    return 0;
}
```

What has changed again?

```
diff -u journal{3,4}.c
```

```
--- journal3.c 2024-02-07 12:31:13.196788801 +0000
+++ journal4.c 2024-02-07 12:31:13.293455473 +0000
@@ -1,6 +1,7 @@
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
+#include <errno.h>

int main(int argc, char *argv[]) {
    char *message = NULL;
@@ -11,6 +12,10 @@

    if (argc < 2) { printf("Usage: %s path/to/log\n", argv[0]); exit(1); };
    FILE *file = fopen(argv[1], "a+");
+   if (file == NULL) {
+       perror("Failed to open log");
+       exit(2);
+   }

    printf("Type your log: ");
    getline(&message, &message_len, stdin);
```

Now when we run...

```
$ ./journal4 <<<hello
Usage ./journal4 path/to/log

$ ./journal4 documents/log.txt <<<hello
Failed to open log: No such file or directory

$ ./journal4 /etc/passwd <<<hello
Failed to open log: Permission denied

$ ./journal4 /dev/stdout
Type your log: hello
20: hello
```

From man 3 strftime:

%c	The preferred date and time representation for the current locale. (The specific format used in the current locale can be obtained by calling <code>nl_langinfo(3)</code> with <code>D_T_FMT</code> as an argument for the <code>%c</code> conversion specification, and with <code>ERA_D_T_FMT</code> for the <code>%Ec</code> conversion specification.) (In the POSIX locale this is equivalent to <code>%a %b %e %H:%M:%S %Y.</code>)
%C	The century number (year/100) as a 2-digit integer. (SU) (The <code>%EC</code> conversion specification corresponds to the name of the era.) (Calculated from <code>tm_year.</code>)

Debugging tools can't catch poorly written code!

But other tools can catch things...

Thinking back to when we fixed up getline... it said it would allocate the memory for the line

► ...did we ever free it?

```
valgrind ./journal4 /dev/stdout <<<hello
```

```
==36111== Memcheck, a memory error detector
==36111== Copyright (C) 2002-2022, and GNU GPL'd, by Julian Seward et al.
==36111== Using Valgrind-3.20.0 and LibVEX; rerun with -h for copyright info
==36111== Command: ./journal4 /dev/stdout
==36111==
20: hello
```

```
Type your log: ==36111==
==36111== HEAP SUMMARY:
==36111== in use at exit: 592 bytes in 2 blocks
==36111== total heap usage: 13 allocs, 11 frees, 13,684 bytes allocated
==36111==
==36111== LEAK SUMMARY:
==36111== definitely lost: 120 bytes in 1 blocks
==36111== indirectly lost: 0 bytes in 0 blocks
==36111== possibly lost: 0 bytes in 0 blocks
==36111== still reachable: 472 bytes in 1 blocks
==36111==    suppressed: 0 bytes in 0 blocks
==36111== Rerun with --leak-check=full to see details of leaked memory
==36111==
==36111== For lists of detected and suppressed errors, rerun with: -s
==36111== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

Wrap up

In this lecture we've gone over the *very basics* of several debugging tools

- ▶ *strace, ltrace, valgrind and gdb will help deal with most of the bugs you encounter*

But so will good defensive programming strategies

- ▶ *Always check the return code of functions*
- ▶ *Always check assumptions*
- ▶ *Always fix your compiler warnings*

...actually get more warnings!

Compiling with the `-Wall -Wextra --std=c11 -pedantic` will make the compiler really picky about your *C* code...

But there are *other tools* called *linters* that can get even more picky

C/C++ Clang Static Analyser, Rats

Java FindBugs

Haskell hlint

Python pylint, mypy


Shellscript shellcheck

Other tools for C/C++ can add extra runtime checks

ASan Address Sanitizer; checks for pointer shenangians

UBSan Undefined Behaviour Sanitizer; checks for *C* gotchas

BPF tools



Linux has a (reasonably) new instrumentation framework called *eBPF*

- ▶ It lets you get *loads* of detail about what programs are doing
- ▶ Highly Linux specific
- ▶ I need to learn it :-)

This weeks lab

Is brand new!

- ▶ I'm gonna give you 5 crackmes
- ▶ They'll ask you for the password
 - ▶ You have to work out what it is

Practice using the debugging tools to work out *what the program is expecting*.

Hmmm...

We're early again...

- ▶ Last year we tried to cram too much stuff into too few lectures...
- ▶ This year I seem to have gone too far the other way (sorry about that)

Well we could show you another use for debugging tools

- ▶ Sometimes being able to see what a program is doing is useful for **other** things ;-)